



The Art of SOCRATIC QUESTIONING: Recursive Thinking with Large Language Models

Jingyuan Qi^{*♦} Zhiyang Xu^{*♦} Ying Shen^{†♣} Minqian Liu^{†♣}

Di Jin[♡] Qifan Wang[♣] Lifu Huang[♣]

[♦]Virginia Tech [♡]Amazon Inc. [♣]Meta AI

{jingyq1,zhiyangx,yings,minqianliu,lifuh}@vt.edu

djinamzn@amazon.com wqfcr@meta.com

Abstract

Chain-of-Thought (CoT) prompting enables large language models to solve complex reasoning problems by generating intermediate steps. However, confined by its inherent single-pass and sequential generation process, CoT heavily relies on the initial decisions, causing errors in early steps to accumulate and impact the final answers. In contrast, humans adopt recursive thinking when tackling complex reasoning problems, i.e., iteratively breaking the original problem into approachable sub-problems and aggregating their answers to resolve the original one. Inspired by the human cognitive process, we propose SOCRATIC QUESTIONING, a divide-and-conquer style algorithm that mimics the recursive thinking process. Specifically, SOCRATIC QUESTIONING leverages large language models to raise and answer sub-questions until collecting enough information to tackle the original question. Unlike CoT, SOCRATIC QUESTIONING explicitly navigates the thinking space, stimulates effective recursive thinking, and is more robust towards errors in the thinking process. Extensive experiments on several complex reasoning tasks, including MMLU, MATH, LogiQA, and visual question-answering demonstrate significant performance improvements over the state-of-the-art prompting methods, such as CoT, and Tree-of-Thought. The qualitative analysis clearly shows that the intermediate reasoning steps elicited by SOCRATIC QUESTIONING are similar to humans' recursively thinking process of complex reasoning problems¹².

1 Introduction

The art of Socratic Questioning is important for critical thinkers and excellence of thought. What Socratic adds is systematicity, depth, and a keen

^{1*} Co-first Authors, [†] Co-second Authors

²All the programs and necessary resources are released in <https://github.com/VT-NLP/SOCRATIC-QUESTIONING>

?

Question: A ball slides down a landslide of height h , the length of the landslide is l , ball mass is m , and the angle between the landslide and the ground is θ . Find the velocity v of the ball at the landslide bottom. Ignore friction force and air resistance.

Standard Prompting

X $v = \sqrt{2 * g * h * \sin(\theta) / (1 + (l * \sin(\theta)) / (2 * h))}$

Chain-of-Thought (CoT)

The kinetic energy theorem states that the net work done on an object equals its change in kinetic energy. In this case, the net work done is done by the frictional force. The work done by friction can be calculated as $W = f * l$, where l is the landslide length and $f = \mu * N$. Write the equation as: $\mu * m * g * \cos(\theta) * l = (1/2) * m * v^2$. Solve that equation, $v = \sqrt{(2 * \mu * g * \cos(\theta) * l) / 1}$.

Socratic Questioning

I do not know.

- ?
1. Does this problem obey the energy conservation law?
 2. What is the mechanical energy of the initial state?
 3. What is the mechanical energy in the final state?

1. Because there is no external force doing work, so this question follows energy conservation law.
2. At beginning, ball's velocity is 0, so its kinetic energy is 0. And the potential energy is $mg * h$.
3. At the end, ball's height is 0, so the potential energy is 0. And kinetic energy is $1/2 * m * v^2$.

Again, find the velocity V of the ball at the bottom of the landslide. Ignore friction force.

Solve $0 + mg * h = 1/2 * m * v^2 + 0$, we get ✓ $v = \sqrt{2 * g * h}$.

Figure 1: Example of a complex question solved by the Standard Prompting, Chain-of-Thought, and SOCRATIC QUESTIONING. Accumulated incorrect reasoning are highlighted in red.

interest in assessing the plausibility of things.

- L. ELDER and R. PAUL, 1998

One unique capability that allows humans to excel at solving complex reasoning problems is *recursive thinking*. If the answer is not immediately achievable, humans think deeper by recursively decomposing the complex problem into simpler and solvable sub-problems.

Recently, by scaling up the parameters, large-




Figure 2: **Schematic comparison of various prompting methods.** Each blue rectangle box represents a *thought* serving as an intermediate reasoning step in the problem-solving process. SOCRATIC QUESTIONING incorporates both a **top-down exploration** process (in red line) to deconstruct complex problems into smaller sub-questions and a **bottom-up backtracking** process (in green line) to recursively solve these sub-questions and gather solutions for higher-level problems.

scale language models (LLMs) (Brown et al., 2020; Chung et al., 2022; OpenAI, 2022; Touvron et al., 2023) gain emerging capabilities, such as Chain-of-Thought (CoT) (Wei et al., 2022) which decomposes the complex problem and solves it step by step. Though CoT has been proven to be effective on various complex reasoning tasks, it’s in nature a single-pass and sequential thinking process that generates the next step based on previous steps, thus only exploring a single way of thinking to approach a problem and easily accumulating errors from previous steps (Turpin et al., 2023). In addition, CoT lacks the ability to refine the already generated reasoning path, as shown in Figure 1.

Inspired by the recursive thinking of humans, we propose SOCRATIC QUESTIONING, a novel divide-and-conquer fashion algorithm that prompts language models to solve complex reasoning problems. As shown in Figure 2 (e), SOCRATIC QUESTIONING consists of a **top-down** exploration process and a **bottom-up** backtracking process. Specifically, in the top-down exploration process, the original complex problem is decomposed into simpler or related sub-problems until the sub-problems can be solved. In the bottom-up backtracking process, the solutions to the sub-problems are returned and selectively used to solve the original problem. The fundamental component that drives SOCRATIC QUESTIONING is a **SELF-QUESTIONING** (SQ) module, that leverages large-scale language models to proactively raise and answer questions that are essential to solving the target question. SOCRATIC QUESTIONING recursively backtracks and tailors the intermediate thoughts acquired from SELF-QUESTIONING until reaching an answer to the original input question. It explicitly navigates the thinking space and is more robust towards thinking errors compared with pre-

vious prompting methods including CoT, Self-Consistency Chain-of-Thought (Wang et al., 2023), and Tree-of-Thought (Yao et al., 2023), as shown in Figure 2.

To show the effectiveness of SOCRATIC QUESTIONING, we conduct extensive experiments on various complex reasoning tasks including the chemistry and physics tasks (Hendrycks et al., 2020), mathematical tasks (Hendrycks et al., 2021), and reading comprehension tasks (Liu et al., 2020). Additionally, we showcase the generalizability of our method by conducting experiments with few-shot multimodal reasoning on VQA-V2 (Goyal et al., 2017), OK-VQA (Marino et al., 2019), and AOK-VQA (Schwenk et al., 2022) datasets. Experimental results indicate that SOCRATIC QUESTIONING substantially improves performance over CoT, SC-CoT, and ToT across all language tasks and outperforms several strong baselines in few-shot multimodal reasoning. The qualitative analysis further demonstrates that SOCRATIC QUESTIONING is capable of eliciting the intermediate reasoning steps through SELF-QUESTIONING, like a critical thinker, and solving complex reasoning problems. The main contributions of our paper are as follows:

- We propose SOCRATIC QUESTIONING, a novel prompting algorithm that can navigate the cognitive thinking space in a ***recursive*** manner.
- We introduce the **SELF-QUESTIONING** module, a core component that actively probes complex problems from various perspectives by raising and addressing questions essential for solving the main problem.
- Our approach achieves significant improvements over the previous prompting methods in various complex reasoning tasks.

2 Related Work

Prompting Large Language Models With the scaling of both modal size and corpus size, large language models (LLMs) such as GPT-3 (Brown et al., 2020) and ChatGPT (OpenAI, 2022) have exhibited emergent abilities, including prompting (Brown et al., 2020), in-context learning (Dong et al., 2023), and commonsense reasoning (Wei et al.). One notable example of emergent abilities is the Chain-of-Thought (CoT) (Wei et al., 2022) which steers large language models to resolve complex problems by guiding them to produce a sequence of intermediate steps before giving the final answer. Self-Consistency Chain-of-Thought (SC-CoT) (Wang et al., 2023) improves naive CoT by sampling multiple reasoning paths and selecting the most consistent answer. SC-CoT is based on the assumption that given a complex reasoning problem, multiple reasoning paths can lead to the unique correct answer. Tree-of-Thought (ToT) (Yao et al., 2023) proposes to break the thinking process into small steps and at each step, the language model deliberately decides a set of next steps to try.

Multimodal Reasoning with Large Language Models Recent studies have explored the collaboration among diverse language and visual models (Yang et al., 2022; Zeng et al., 2022; Huang et al., 2022). For example, PICa (Yang et al., 2022) utilize image captions as the bridge between visual model and GPT-3 to perform few-shot knowledge-based VQA. Socratic models (Zeng et al., 2022) present a modular framework that utilizes language-based exchange between pre-trained models and other modules. However, these studies only rely on text as the shared interface, which can inevitably lead to information loss when translating visual information into language. In addition, several concurrent studies (Wu et al., 2023; Surís et al., 2023; Lu et al., 2023) have also explored the utilization of large language models for composing various language and visual models.

Question Decomposition Recent research has underscored the effectiveness of question decomposition and sub-question generation techniques in tackling complex tasks. DECOMPRC (Min et al., 2019), for instance, utilizes a limited amount of human-labeled data to train a span-based sub-question generator and simplifies multi-hop questions into single-hop questions. Similarly, (Nogueira and Cho, 2017) leverages reinforcement

learning for weakly supervised question generation and (Perez et al., 2020) introduces ONUS, an algorithm that harnesses large-scale questions sourced from the internet to perform unsupervised question decomposition. More recently, (Patel et al., 2022) proposes an alternative approach to enhance the performance of LLMs by decomposing challenging questions into simpler sub-questions on various tasks. Notably, the efficacy of question decomposition has been demonstrated across a range of tasks and domains, including solving mathematical problems (Shridhar et al., 2022), medical question answering (Roberts et al., 2014), and factual correction (Huang et al., 2023).

3 Method

3.1 SOCRATIC QUESTIONING

Figure 3 shows the overview of the SOCRATIC QUESTIONING approach, which is essentially a recursive thinking process involving a *top-down exploration process* (in red line) and a *bottom-up backtracking process* (in green line). The top-down exploration process proactively breaks down the question into simpler sub-questions until the sub-questions are answered with high confidence. The bottom-up backtracking process recursively solves questions in which the answers to sub-questions are collected to solve the higher-level more complex questions.

In the beginning, we are given a target question $Q_1^{0,0}$, the context C (if provided), and an optional hint $H_1^{0,0}$. The hint is initially Null but will be updated and enriched as the recursive thinking process continues and results from sub-questions are aggregated. We first run the top-down process to explore the thinking space by invoking the SELF-QUESTIONING module. We use depth d and turn t to identify the node in our reasoning tree. Depth d refers to the traditional depth of the recursion algorithm. Turn t refers to the times of SOCRATIC QUESTIONING invoking the SELF-QUESTIONING module for each question. For example, at depth d , turn t , SELF-QUESTIONING takes in the i_{th} question $Q_i^{d,t}$, hint $H_i^{d,t}$, the context C , and decides if it can answer the question $Q_i^{d,t}$: (1) If SELF-QUESTIONING can directly output the answer $A_i^{d,t}$ for the question $Q_i^{d,t}$ with high confidence, the bottom-up backtracking process starts by converting the answer $A_i^{d,t}$ to a hint $\tilde{H}_i^{d,t}$ with a QA-to-Hint module ($\tilde{H}_i^{0,t}$ equals $A_i^{0,t}$ directly when $d = 0$).




Figure 3: Overview of our SOCRATIC QUESTIONING algorithm.

0) and adding $\tilde{H}_i^{d,t}$ into the hints $H^{d-1,t}$ of the parent question Q^{d-1} . (2) If SELF-QUESTIONING cannot directly output an answer with high confidence, it outputs a set of sub-questions $\mathcal{Q}^{d+1,t}$ related to $Q_i^{d,t}$. Then we run SELF-QUESTIONING on each newly generated sub-question $Q_j^{d+1,t}$ until it's answered with high confidence. Once we obtain the answers to all the sub-questions $\mathcal{Q}^{d+1,t}$, we convert the answers into hints and incorporate them to update $H_i^{d,t}$ to $H_i^{d,t+1}$. We then run SELF-QUESTIONING on $Q_i^{d,t+1}$ again with updated hints $H_i^{d,t+1}$. This recursive process continues until we reach the tree's root and the original question Q_1^0 is answered by \tilde{H}_1^0 . We provide the pseudo-code of SOCRATIC QUESTIONING in Algorithm 1.

3.2 SELF-QUESTIONING

SELF-QUESTIONING is designed to answer the given question, self-check the answer, and raise sub-questions. At depth d , turn t , SELF-QUESTIONING takes in the i th question $Q_i^{d,t}$, the context C (if available), and hints $H_i^{d,t}$ (if available) and tries to generate an answer or a set of related sub-questions. SELF-QUESTIONING consists of two modules, a **Question-Answering (QA) Module** that outputs an answer $A_i^{d,t}$ for $Q_i^{d,t}$ based on C and $H_i^{d,t}$, and an associated confidence level: high, medium, or low. If the confidence of the answer is high, or either depth d or turn t met the pre-defined limit d_m and t_m , SELF-QUESTIONING invokes the QA2H module to merge the question $Q_i^{d,t}$ and answer $A_i^{d,t}$ to hint $\tilde{H}_i^{d,t}$ as output (when $d = 0$, we skip the merging process because the

Algorithm 1: SOCRATIC QUESTIONING

Input: Question $Q_i^{d,t}$, Hint $H_i^{d,t}$, Context C , Current Depth d , Max Depth d_m , Current Turn t , Max Turn t_m , Question Answer Prompt P_{QA} , Question Generate Prompt P_{QG} , QA to Hint Prompt P_{QA2H}

Output: Hint $\tilde{H}_i^{d,t}$

```

1 for  $t \leq t_m$  do
2   // call self-questioning
3    $< Q^{d+1,t}, H^{d+1,t}, C > \leftarrow$ 
4   SELF-QUESTIONING( $Q_i^{d,t}, H_i^{d,t}, C,$ 
5    $d, d_m, t, t_m, P_{QA}, P_{QG}$ );
6   if  $Q^{d+1,t} \neq \emptyset$  then
7     for each  $Q_j^{d+1,t} \in Q^{d+1,t}$  do
8       // recursively answer
9       sub-questions
10       $\tilde{H}_j^{d+1,t} \leftarrow$ 
11      SOCRATIC QUESTIONING( $Q_j^{d+1,t},$ 
12       $H_j^{d+1,t}, C, d + 1, d_m, t, t_m, P_{QA},$ 
13       $P_{QG}$ );
14      // gather hint
15       $H^d.insert(\tilde{H}_j^{d+1,t});$ 
16    else
17       $\tilde{H}_j^{d,t} \leftarrow H^{d+1,t}[0];$ 
18      return  $\tilde{H}_j^{d,t};$ 
19     $t \leftarrow t + 1;$ 

```

answer A_1^0 , is the final answer and does not need to be rewritten to hint). Both Max Depth d_m and Max Turn t_m prevent SOCRATIC QUESTIONING from infinite recursion. On the other hand, if the confidence of the answer is lower than high, a **Question-Generation (QG) Module** is called to generate a set of sub-questions $\{Q_0^{d+1,t}, \dots, Q_n^{d+1,t}\}$ to collect more information based on $Q_i^{d,t}$, C , and

Algorithm 2: SELF-QUESTIONING

Input: Question $Q_i^{d,t}$, Hint $H_i^{d,t}$, Context C , Current Depth d , Max Depth d_m , Current Turn t , Max Turn t_m , Question Answer Prompt P_{QA} , Question Generate Prompt P_{QG} , QA to Hint Prompt P_{QA2H}

Output: $\langle Q^{d+1,t}, \mathcal{H}^{d+1,t}, C \rangle$

```

1 Must_Answer ← False;
2 if  $d = d_m$  or  $t = t_m$  then
3   | Must_Answer ← True;
  // call the Question-Answering module
4  $\langle A_i^{d,t}, confidence \rangle \leftarrow$ 
  QA( $Q_i^{d,t}, H_i^{d,t}, C, P_{QA}$ );
5 if confidence = high or Must_Answer then
6   | if  $d \neq 0$  then
     |   // merge QA to a hint
     |    $\tilde{H}_i^{d,t} \leftarrow QA2H(Q_i^{d,t}, A_i^{d,t}, P_{QA2H})$ ;
8   | else
9     |    $\tilde{H}_i^{d,t} \leftarrow A^d$ ;
10  |  $\mathcal{Q}^{d+1} \leftarrow \emptyset$ ;
11  |  $\mathcal{H}^{d+1,t} \leftarrow \{\tilde{H}_i^{d,t}\}$ ;
12 else
  // call the Question-Generation module
13  |  $\mathcal{Q}^{d+1,t} \leftarrow QG(Q_i^{d,t}, H_i^{d,t}, C, P_{QG})$ ;
14  |  $\mathcal{H}^{d+1,t} \leftarrow \emptyset$ ;
15 return  $\langle Q^{d+1,t}, \mathcal{H}^{d+1,t}, C \rangle$ ;
```

$H_i^{d,t}$, where $n < n_m$ and n_m denotes the maximum number of sub-questions to be generated. Algorithm 2 shows the pseudo-code of the SELF-QUESTIONING algorithm.

3.2.1 Question-Answering (QA) Module

The QA module aims to answer either the target question or a sub-question asked by the SELF-QUESTIONING module, based on the optional context and hints. We propose to leverage a large-scale language model (LLM), such as GPT-3 or ChatGPT (OpenAI, 2022), to answer the question given their superior reasoning capabilities demonstrated in previous studies (Brown et al., 2020; Zhang et al., 2022; Wei et al., 2022; Touvron et al., 2023; Yao et al., 2023).

Specifically, the input to the QA module consists of the given question $Q_i^{d,t}$, the context C , the optional hints $H_i^{d,t}$, and a prompt P_{QA} designed to guide the QA module to generate an answer $A_i^{d,t}$ based on the inputs and output a confidence level. When the hints $H_i^{d,t}$ are available, P_{QA} also asks the QA module to indicate which hints are used to produce the answer.

$$A_i^{d,t}, confidence = QA(Q_i^{d,t}, H_i^{d,t}, C, P_{QA}), \quad (1)$$

where $confidence \in \{high, medium, low\}$.

3.2.2 Question-Generation (QG) Module

When the QA module outputs an answer for question $Q_i^{d,t}$ with low confidence, it's very likely that the answer is not correct and we need to collect additional hints to help the QA module produce a more confident answer. To do so, we design a Question-Generation (QG) module to raise a set of sub-questions that are related to $Q_i^{d,t}$. The QG module is also based on a large language model, such as ChatGPT, that takes the question $Q_i^{d,t}$, optional hints $H_i^{d,t}$, the context C , and a prompt P_{QG} as input and outputs a set of sub-questions:

$$\{\mathcal{Q}_0^{d+1}, \dots, \mathcal{Q}_n^{d+1}\} = QG(Q_i^{d,t}, H_i^{d,t}, C, P_{QG}), \quad (2)$$

where $n < n_m$. Intuitively, the sub-questions should be simpler than $Q_i^{d,t}$ and more likely to be answered by the QA module with high confidence.

3.2.3 QA-to-Hint (QA2H) Module

Since the answers to sub-questions may not be self-contained, we further design a QA-to-Hint module (QA2H) to merge each sub-question with its answer into a statement. Specifically, we feed the sub-question $Q_i^{d,t}$ and its answer $A_i^{d,t}$ to an LLM with the prompt P_{QA2H} which asks the LLM to rewrite the question to a statement by incorporating the answer:

$$\tilde{H}^d = QA2H(Q_i^{d,t}, A_i^{d,t}, P_{QA2H}), \quad (3)$$

4 SOCRATIC QUESTIONING for Few-Shot Multimodal Reasoning

SOCRATIC QUESTIONING can be naturally applied to text-based complex reasoning tasks as all the key components are based on large language models, such as ChatGPT. There are two critical challenges when applying SOCRATIC QUESTIONING to multimodal reasoning: (1) the language model cannot process visual information, and (2) simply applying a generic captioning model to convert visual content to natural language may not capture the key information required to answer a question.

Converting Visual Information into Context

We propose to leverage LLMs to answer visual questions since some of the visual questions are knowledge-demanding (Marino et al., 2019; Schwenk et al., 2022) and LLMs are capable of storing commonsense knowledge and excel in complex reasoning tasks (Brown et al., 2020; Wei et al., 2022; Wang et al., 2023). To overcome the LLMs'

shortcomings that they cannot perceive visual information, previous works (Yang et al., 2022; Zeng et al., 2022) leverage an image captioning model to convert visual information into text and use LLMs to perform few-shot visual question answering (VQA) tasks. However, considering the richness and density of the information contained in an image, a generic caption may not be able to capture the key information that is necessary to answer a question. Thus, in order to adapt our SOCRATIC QUESTIONING, we employ a visual perception model, BLIP-2 (Li et al., 2023), to describe the content of the image that is specific to a prompt. The input to BLIP-2 is an image I (i.e., the image input of the VQA task) and a text prompt Q , and the output is an image caption C describing the part of the image related to the prompt: $C = \text{BLIP-2}(I, Q)$, where the text prompt Q corresponds to Q^d in Equation (1) and the caption C corresponds to the context C in Equation (1). By leveraging the visual perception model, we are able to resolve the hindrance and adopt our SOCRATIC QUESTIONING framework on VQA. We show more details on how we adapt SOCRATIC QUESTIONING to VQA in Appendix A.

5 Experiment Setups

Language-Only Tasks We leverage ChatGPT as the LLM for QA, QG, and QA2H modules, and provide detailed prompts for each module in Appendix K. We evaluate SOCRATIC QUESTIONING on several complex reasoning tasks, including the Physics and Chemistry tasks in **Massive Multitask Language Understanding (MMLU)** (Hendrycks et al., 2020), Mathematical tasks in MATH (Hendrycks et al., 2021), and logical reasoning tasks based on LogiQA (Liu et al., 2020). We adopt several state-of-the-art prompting methods as baselines, including **Standard Prompting (SP)** that directly prompts ChatGPT to answers a question with a few in-context examples. **Chain-of-Thought (CoT)** (Wei et al., 2022), **Self-Consistency Chain-of-Thought (SC-CoT)** (Wang et al., 2023), and **Tree-of-Thought (ToT)** (Yao et al., 2023). Following previous studies (Chowdhery et al., 2023; Hoffmann et al., 2022), we use exact match to measure the accuracy for all language-only tasks. More details for the baselines, evaluation metrics, and evaluation datasets are discussed in Appendix C.1.

Multimodal Tasks We use blip2-flan-t5-xl as our Visual Perception module. We leverage ChatGPT (OpenAI, 2022) for Factual/Visual Question Generation and Factual Question Answering and GPT-3 (GPT-3-davinci-003) for Visual Question Answering³, motivated by the observation that ChatGPT tends to be excessively cautious and neutral, and avoids answering some questions. We provide detailed sample prompts for each module in Appendix K. We evaluate SOCRATIC QUESTIONING on several visual question answering datasets, including **VQA-V2** (Goyal et al., 2017), **OK-VQA** (Marino et al., 2019) and **AOK-VQA** (Schwenk et al., 2022), and compare our approach with several baselines, including **BLIP-2** (Li et al., 2023) and **PICa** (Yang et al., 2022). More details for implementation, baselines, and datasets are discussed in Appendix C.2. For evaluation, we employ the conventional VQA accuracy metric (Goyal et al., 2017) to measure the performance. To alleviate stringent penalization for minor discrepancies between predicted answers and ground truth, we normalize the answers by converting plural forms to singular forms and changing the tense of verbs to present tense. In addition, to address the conventional metric’s limitation due to synonyms and expression differences, we design semantic-based accuracy by employing ChatGPT to evaluate the correctness of the predicted answers (Fu et al., 2023; Liu et al., 2023b). We provide ChatGPT with the visual question, the predicted answer and the ground-truth answer, and ask if the ground-truth answer and the predicted answer can support each other. If the answer is "Yes", we treat the predicted answer as correct. We show the exact prompts used for ChatGPT in Appendix K.8.

6 Results and Discussions

6.1 Quantitative Results

Language-only Tasks Table 1 shows the quantitative results in terms of accuracy for language-only reasoning tasks. Our method substantially outperforms previous state-of-the-art methods by 4.34%, 2.98%, 4.22%, and 4.66% absolute gains in MATH, Physics, Chemistry, and Logic benchmarks, respectively. This effectively demonstrates the superiority of our approach. We also conduct an experiment on how the maximum number of turns t_m affects the performance. Specifically, we experiment with the setting where $t_m = 2$ (2-Turns)

³These components are detailed in Appendix A.

	MATH (DA)	MMLU Physics	MMLU Chemistry	LogiQA	Avg
Standard-Prompting	7.00	65.11	53.20	54.67	45.00
CoT (Wei et al., 2022)	7.33	67.66	57.14	48.33	45.12
SC-CoT (Wei et al., 2022)	7.00	68.51	59.33	49.00	46.03
ToT (Yao et al., 2023)	0.00	40.00	26.60	22.22	29.46
SOCRATIC QUESTIONING (2-Turns)	<u>7.67</u>	71.49	63.55	59.33	<u>50.51</u>
SOCRATIC QUESTIONING (3-Turns)	11.67	69.36	63.55	58.00	50.65

Table 1: Accuracy (%) using Exact Match. The best performance is highlighted in **bold** and the second best performance is highlighted with underline.




Figure 4: Qualitative results of CoT, ToT, and SOCRATIC QUESTIONING on the Physics task. The correct answer of this example is **D**.

Model	VQA-V2	OK-VQA	AOK-VQA
BLIP-2 (Li et al., 2023)	36.7	21.14	0
PICa (Yang et al., 2022)	<u>43.18</u>	<u>29.94</u>	<u>28.6</u>
SOCRATIC QUESTIONING	46.64	31.24	29.58

Table 2: Traditional VQA Accuracy (%) based on Exactly Match. The best performance is highlighted in **bold** and the second best performance is highlighted with underline.

and $t_m = 3$ (3-Turns). From Table 1, the model with maximum 2 turns achieves better performance on Physics and LogiQA datasets, while the model with $t_m = 3$ performs better on the MATH dataset. One possible reason is that the Physics and LogiQA benchmarks may not be challenging enough and reasoning within 2 turns is sufficient to answer most of the questions. We provide a concrete example in Appendix G.1.

Multimodal Tasks Table 2 and 3 show the quantitative results using traditional VQA accuracy and semantic-based accuracy, respectively. For both results, our SOCRATIC QUESTIONING method outperforms the previous state-of-the-art approaches

Model	VQA-V2	OK-VQA	AOK-VQA
BLIP-2 (Li et al., 2023)	57.2	46.75	43.29
PICa (Yang et al., 2022)	49.8	<u>48.05</u>	<u>46.85</u>
SOCRATIC QUESTIONING	<u>54.4</u>	53.03	49.55

Table 3: Semantic-based VQA Accuracy (%) using NLI. The best performance is highlighted in **bold** and the second best performance is highlighted with underline.

on most benchmarks, often by a large margin. The only exception is semantic-based accuracy on the VQA-V2 dataset. A possible reason is that the tasks on VQA-V2 focus more on the visual recognition and detection aspect and do not require much reasoning capability and external knowledge.

6.2 Qualitative Result

Language-only Tasks Figure 4 shows the qualitative results of SOCRATIC QUESTIONING and baselines on the Physics task. As one can observe, SOCRATIC QUESTIONING can effectively prompt hints containing the necessary information to solve the original problem and selectively use the hints to reach the correct final answer. On the other hand, CoT and ToT reach the wrong answer due to the




Figure 5: Qualitative results of few-shot VQA using BLIP-2, PICa, and SOCRATIC QUESTIONING (2-Depth 2-Turn).

	Answered Correctly	Answered Incorrectly
Avg. Hints	3.28	3.68
Avg. Depth	2.89	2.92

Table 4: Averaged numbers of hints and depth of SOCRATIC QUESTIONING used for questions answered correctly and incorrectly, respectively.

poorly sampled reasoning path.

Multimodal Tasks Figure 5 shows several examples of few-shot VQA tasks from the baselines and SOCRATIC QUESTIONING. We demonstrate that the hints acquired via the sub-problems are highly related to the original problem (e.g., “*weather conditions are cold*”), and by considering the collected hints, the SOCRATIC QUESTIONING reaches the correct final answer (e.g., “*warmth*”). In contrast, the answer from BLIP-2 is irrelevant to the given question, due to the generic caption.

6.3 How do the Numbers of Turns and Depths Affect the Model?

Performance Breakdown w/ Number of Turns To study how the number of reasoning turns affects the performance across different benchmarks, we investigate how the baselines and our method per-




Figure 6: Accuracy (%) on the examples that triggered 2 turns of reasoning by SOCRATIC QUESTIONING.




Figure 7: Accuracy (%) on the examples that triggered 3 turns of reasoning by SOCRATIC QUESTIONING.

form on the examples that triggered 2 and 3 turns of reasoning by SOCRATIC QUESTIONING in Figure 6 and Figure 7, respectively. This experiment can be considered as breaking down the results in Table 1 into two groups based on the number of reasoning turns. From Figure 6, our approach outperforms the baselines on all benchmarks except for the MATH dataset. From Figure 7, our approach outperforms the baselines on relatively challenging tasks such as MATH but performs more poorly on easier tasks such as Physics. This indicates SOCRATIC QUESTIONING with more turns can tackle challenging problems more effectively.

The Effect of Hyperparameters t_m and d_m In addition to the discussion in 6.1, we conduct a more in-depth analysis of how the maximum number of turns t_m and maximum number of depths d_m affect the performance of our SOCRATIC QUESTIONING. In Figure 8, we show the heat map under different hyperparameter settings, where the number in each cell is the accuracy (%) given a specific combination of t_m and d_m . We observe two general trends: (1) the accuracy increases when t_m gets larger, and




Figure 8: Quantitative results SOCRATIC QUESTIONING on the MATH dataset with different values of the hyperparameters t_m and d_m .

(2) the accuracy decreases when d_m gets larger. These results imply that our approach can benefit from raising more questions directly related to the original question. Also, performing reasoning with a larger maximum depth does not yield better performance since the benchmark may not be challenging enough, and exploring at a deeper level may introduce irrelevant information. We provide a concrete example in Appendix G.2. In addition, we analyze the computational cost of SOCRATIC QUESTIONING compared to other baselines in Appendix H, and show that while achieving stronger performance, our proposed algorithm enjoys higher efficiency than most of baselines.

6.4 How does the Difficulty of Questions Affect the Model?

Table 4 presents the averaged numbers of hints and depth used to answer the original questions for correct and incorrect answers. As one can observe, for incorrect answers, the LLM raises more sub-questions, which demonstrates that the LLM tends to explore more thinking space when tackling questions that it does not know the answers. This trend also agrees with the depth. If the question is hard for the LLM, the model tends to break the sub-questions into even more basic questions.

7 Conclusion

We present SOCRATIC QUESTIONING, a novel divide-and-conquer fashion algorithm that is inspired by human’s recursive thinking processes. SOCRATIC QUESTIONING consists of a top-down reasoning phase that decomposes a complex problem into simpler sub-problems and a bottom-top phase where the solutions to the sub-problems are recursively returned and used to solve the original problem at higher levels. Extensive experiments

on four challenging language-only tasks and the few-shot VQA task validate the effectiveness of our SOCRATIC QUESTIONING. Moreover, qualitative analysis demonstrates our approach can effectively elicit intermediate reasoning steps and consequently yield a correct final answer while enjoying transparency and interpretability.

Limitation

The self-checking functionality lacks sufficient sensitivity to incorrect responses, as its confidence estimation heavily relies on LLMs themselves. While we employed ChatGPT as the backbone for our algorithm, its tendency towards overconfidence leads to a low frequency of sub-question generation.

Our study exhibits a lack of diversity in visual models used to extract information from images. We only use BLIP-2 (Li et al., 2023) as an image caption model in current experiments. However, the incorporation of diverse visual models, such as dense caption models, Optical Character Recognition (OCR), or scene graph models, may potentially yield a broader spectrum of image information, thus facilitating the resolution of sub-questions. In addition, to help BLIP-2 to better follow instructions from LLMs, we propose to leverage recent techniques developed in visual instruction tuning (Liu et al., 2023a; Xu et al., 2023b,a; Dai et al., 2023).

Additionally, our experiments were constrained to the English language datasets and we only consider the VQA task to showcase the multi-modal performance. However, given the generality of our algorithm, we plan to test its functionality with multilingual datasets and experiment it on other domains, such as speech (You et al., 2020, 2022), and video (Rose et al., 2023).

Acknowledgments

This research is based upon work supported by the U.S. DARPA ECOLE Program # HR001122S0052. The views and conclusions contained herein are those of the authors and should not be interpreted as necessarily representing the official policies, either expressed or implied, of DARPA or the U.S. Government. The U.S. Government is authorized to reproduce and distribute reprints for governmental purposes notwithstanding any copyright annotation therein.

References

- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901.
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, Parker Schuh, Kensen Shi, Sasha Tsvyashchenko, Joshua Maynez, Abhishek Rao, Parker Barnes, Yi Tay, Noam Shazeer, Vinodkumar Prabhakaran, Emily Reif, Nan Du, Ben Hutchinson, Reiner Pope, James Bradbury, Jacob Austin, Michael Isard, Guy Gur-Ari, Pengcheng Yin, Toju Duke, Anselm Levskaya, Sanjay Ghemawat, Sunipa Dev, Henryk Michalewski, Xavier Garcia, Vedant Misra, Kevin Robinson, Liam Fedus, Denny Zhou, Daphne Ippolito, David Luan, Hyeontaek Lim, Barret Zoph, Alexander Spiridonov, Ryan Sepassi, David Dohan, Shivani Agrawal, Mark Omernick, Andrew M. Dai, Thanumalayan Sankaranarayana Pillai, Marie Pellat, Aitor Lewkowycz, Erica Moreira, Rewon Child, Oleksandr Polozov, Katherine Lee, Zongwei Zhou, Xuezhi Wang, Brennan Saeta, Mark Diaz, Orhan Firat, Michele Catasta, Jason Wei, Kathy Meier-Hellstern, Douglas Eck, Jeff Dean, Slav Petrov, and Noah Fiedel. 2023. **Palm: Scaling language modeling with pathways**. *J. Mach. Learn. Res.*, 24:240:1–240:113.
- Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Eric Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, Albert Webson, Shixiang Shane Gu, Zhuyun Dai, Mirac Suzgun, Xinyun Chen, Aakanksha Chowdhery, Sharan Narang, Gaurav Mishra, Adams Yu, Vincent Y. Zhao, Yanping Huang, Andrew M. Dai, Hongkun Yu, Slav Petrov, Ed H. Chi, Jeff Dean, Jacob Devlin, Adam Roberts, Denny Zhou, Quoc V. Le, and Jason Wei. 2022. **Scaling instruction-finetuned language models**. *CoRR*, abs/2210.11416.
- Wenliang Dai, Junnan Li, Dongxu Li, Anthony Meng Huat Tiong, Junqi Zhao, Weisheng Wang, Boyang Li, Pascale Fung, and Steven C. H. Hoi. 2023. **Instructclip: Towards general-purpose vision-language models with instruction tuning**. *CoRR*, abs/2305.06500.
- Qingxiu Dong, Lei Li, Damai Dai, Ce Zheng, Zhiyong Wu, Baobao Chang, Xu Sun, Jingjing Xu, Lei Li, and Zhifang Sui. 2023. **A survey for in-context learning**. *CoRR*, abs/2301.00234.
- Jinlan Fu, See-Kiong Ng, Zhengbao Jiang, and Pengfei Liu. 2023. **Gptscore: Evaluate as you desire**. *CoRR*, abs/2302.04166.
- Mor Geva, Daniel Khashabi, Elad Segal, Tushar Khot, Dan Roth, and Jonathan Berant. 2021. **Did aristotle use a laptop? a question answering benchmark with implicit reasoning strategies**. *Transactions of the Association for Computational Linguistics*, 9:346–361.
- Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. 2017. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 6904–6913.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2020. **Measuring massive multitask language understanding**. *CoRR*, abs/2009.03300.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. **Measuring mathematical problem solving with the MATH dataset**. *CoRR*, abs/2103.03874.
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al. 2022. Training compute-optimal large language models. *arXiv preprint arXiv:2203.15556*.
- Kung-Hsiang Huang, Hou Pong Chan, and Heng Ji. 2023. **Zero-shot faithful factual error correction**. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5660–5676, Toronto, Canada. Association for Computational Linguistics.
- Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, Yevgen Chebotar, Pierre Sermanet, Tomas Jackson, Noah Brown, Linda Luu, Sergey Levine, Karol Hausman, and Brian Ichter. 2022. **Inner monologue: Embodied reasoning through planning with language models**. In *Conference on Robot Learning, CoRL 2022, 14–18 December 2022, Auckland, New Zealand*, volume 205 of *Proceedings of Machine Learning Research*, pages 1769–1782. PMLR.
- Zhengbao Jiang, Jun Araki, Haibo Ding, and Graham Neubig. 2021. **How can we know when language models know? on the calibration of language models for question answering**. *Transactions of the Association for Computational Linguistics*, 9:962–977.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislav Fort, Deep Ganguli, Danny Hernandez, Josh Jacobson, Jackson Kernion, Shauna Kravec, Liane Lovitt, Kamal Ndousse, Catherine Olsson, Sam Ringer, Dario Amodei, Tom Brown, Jack Clark, Nicholas Joseph, Ben Mann, Sam McCandlish, Chris Olah, and Jared

- Kaplan. 2022. *Language models (mostly) know what they know*. *CoRR*, abs/2207.05221.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven C. H. Hoi. 2023. *BLIP-2: bootstrapping language-image pre-training with frozen image encoders and large language models*. In *International Conference on Machine Learning, ICML 2023, 23-29 July 2023, Honolulu, Hawaii, USA*, volume 202 of *Proceedings of Machine Learning Research*, pages 19730–19742. PMLR.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. 2023a. Visual instruction tuning.
- Jian Liu, Leyang Cui, Hanmeng Liu, Dandan Huang, Yile Wang, and Yue Zhang. 2020. *Logica: A challenge dataset for machine reading comprehension with logical reasoning*. In *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence, IJCAI 2020*, pages 3622–3628. ijcai.org.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023b. *G-eval: NLG evaluation using GPT-4 with better human alignment*. *CoRR*, abs/2303.16634.
- Pan Lu, Baolin Peng, Hao Cheng, Michel Galley, Kai-Wei Chang, Ying Nian Wu, Song-Chun Zhu, and Jianfeng Gao. 2023. *Chameleon: Plug-and-play compositional reasoning with large language models*. *CoRR*, abs/2304.09842.
- Kenneth Marino, Mohammad Rastegari, Ali Farhadi, and Roozbeh Mottaghi. 2019. Ok-vqa: A visual question answering benchmark requiring external knowledge. In *Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Sewon Min, Victor Zhong, Luke Zettlemoyer, and Hanneh Hajishirzi. 2019. *Multi-hop reading comprehension through question decomposition and rescoring*. *CoRR*, abs/1906.02916.
- Rodrigo Nogueira and Kyunghyun Cho. 2017. *Task-oriented query reformulation with reinforcement learning*. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 574–583, Copenhagen, Denmark. Association for Computational Linguistics.
- OpenAI. 2022. ChatGPT: Optimizing language models for dialogue. <https://openai.com/blog/chatgpt/>.
- Pruthvi Patel, Swaroop Mishra, Mihir Parmar, and Chitta Baral. 2022. *Is a question decomposition unit all we need?* In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 4553–4569, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Ethan Perez, Patrick S. H. Lewis, Wen-tau Yih, Kyunghyun Cho, and Douwe Kiela. 2020. *Unsupervised question decomposition for question answering*. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing, EMNLP 2020, Online, November 16-20, 2020*, pages 8864–8880. Association for Computational Linguistics.
- Chengwei Qin, Aston Zhang, Zhuosheng Zhang, Jiaao Chen, Michihiro Yasunaga, and Diyi Yang. 2023. *Is chatgpt a general-purpose natural language processing task solver?*
- Kirk Roberts, Kate Masterton, Marcelo Fiszman, Halil Kilicoglu, and Dina Demner-Fushman. 2014. *Annotating question decomposition on complex medical questions*. In *Proceedings of the Ninth International Conference on Language Resources and Evaluation, LREC 2014, Reykjavik, Iceland, May 26-31, 2014*, pages 2598–2602. European Language Resources Association (ELRA).
- Daniel Rose, Vaishnavi Himakunthala, Andy Ouyang, Ryan He, Alex Mei, Yuje Lu, Michael Saxon, Chinmay Sonar, Diba Mirza, and William Yang Wang. 2023. *Visual chain of thought: Bridging logical gaps with multimodal infillings*. *CoRR*, abs/2305.02317.
- Dustin Schwenk, Apoorv Khandelwal, Christopher Clark, Kenneth Marino, and Roozbeh Mottaghi. 2022. A-okvqa: A benchmark for visual question answering using world knowledge. In *Computer Vision-ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part VIII*, pages 146–162. Springer.
- Kumar Shridhar, Jakub Macina, Mennatallah El-Assady, Tanmay Sinha, Manu Kapur, and Mrinmaya Sachan. 2022. *Automatic generation of socratic subquestions for teaching math word problems*. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 4136–4149. Association for Computational Linguistics.
- Dídac Surís, Sachit Menon, and Carl Vondrick. 2023. *Vipergpt: Visual inference via python execution for reasoning*. *CoRR*, abs/2303.08128.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurélien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. 2023. *Llama: Open and efficient foundation language models*. *CoRR*, abs/2302.13971.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. *Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting*. *CoRR*, abs/2305.04388.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. *Self-consistency improves chain of thought reasoning in language*

- models.** In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net.
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, et al. Emergent abilities of large language models.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Ed H. Chi, Quoc Le, and Denny Zhou. 2022. **Chain of thought prompting elicits reasoning in large language models.** *CoRR*, abs/2201.11903.
- Chenfei Wu, Shengming Yin, Weizhen Qi, Xiaodong Wang, Zecheng Tang, and Nan Duan. 2023. **Visual chatgpt: Talking, drawing and editing with visual foundation models.** *CoRR*, abs/2303.04671.
- Zhiyang Xu, Trevor Ashby, Chao Feng, Rulin Shao, Ying Shen, Di Jin, Qifan Wang, and Lifu Huang. 2023a. **Vision-flan: Scaling visual instruction tuning.**
- Zhiyang Xu, Ying Shen, and Lifu Huang. 2023b. **Multi-instruct: Improving multi-modal zero-shot learning via instruction tuning.** In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2023, Toronto, Canada, July 9-14, 2023*, pages 11445–11465. Association for Computational Linguistics.
- Zhengyuan Yang, Zhe Gan, Jianfeng Wang, Xiaowei Hu, Yumao Lu, Zicheng Liu, and Lijuan Wang. 2022. **An empirical study of GPT-3 for few-shot knowledge-based VQA.** In *Thirty-Sixth AAAI Conference on Artificial Intelligence, AAAI 2022, Thirty-Fourth Conference on Innovative Applications of Artificial Intelligence, IAAI 2022, The Twelfth Symposium on Educational Advances in Artificial Intelligence, EAAI 2022 Virtual Event, February 22 - March 1, 2022*, pages 3081–3089. AAAI Press.
- Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. 2023. **Tree of thoughts: Deliberate problem solving with large language models.** *CoRR*, abs/2305.10601.
- Chenyu You, Nuo Chen, Fenglin Liu, Shen Ge, Xian Wu, and Yuexian Zou. 2022. End-to-end spoken conversational question answering: Task, dataset and model. In *Findings of NAACL 2022*.
- Chenyu You, Nuo Chen, Fenglin Liu, Dongchao Yang, and Yuexian Zou. 2020. Towards data distillation for end-to-end spoken conversational question answering. *arXiv preprint arXiv:2010.08923*.
- Andy Zeng, Adrian Wong, Stefan Welker, Krzysztof Choromanski, Federico Tombari, Aveek Purohit, Michael S. Ryoo, Vikas Sindhwani, Johnny Lee, Vincent Vanhoucke, and Pete Florence. 2022. **Socratic models: Composing zero-shot multimodal reasoning with language.** *CoRR*, abs/2204.00598.
- Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuhui Chen, Christopher Dewan, Mona T. Diab, Xian Li, Xi Victoria Lin, Todor Mihaylov, Myle Ott, Sam Shleifer, Kurt Shuster, Daniel Simig, Punit Singh Koura, Anjali Sridhar, Tianlu Wang, and Luke Zettlemoyer. 2022. **OPT: open pre-trained transformer language models.** *CoRR*, abs/2205.01068.

A Adapting SOCRATIC QUESTIONING to Visual Question Answering

Question-Generation (QG) Module Some tasks (e.g., OK-VQA, AOK-VQA) require commonsense knowledge. Although LLMs can retrieve knowledge from its parameter, they are prone to hallucination and the black-box retrieving process is hard to debug. In order to gain a clear understanding of the factual knowledge used in answering a question, we divide the QG module in Section 3.2.2 into two sub-modules: A Fact-Question-Generation (FQG) sub-module which generates factual questions related to background knowledge of the given question, and a Visual-Question-Generation (VQG) sub-module generates visual questions, which aims to guide the Visual Perception module to focus on question-related image regions and seek more image information.

Question-Answering (QA) Module To accommodate the two question types, we also divide the QA module in section 3.2.1 into two sub-modules: A Factual-Question-Answering module (FQA) and a Visual-Question-Answering module (VQA). Both FQA and VQA modules follow the same formulation in Equation (1). The input C to VQA is the caption related to the question Q^d and is prompted via the Equation of BLIP-2.

SELF-QUESTIONING Figure 9 demonstrates the detailed step of the SELF-QUESTIONING algorithm in the multimodal setting. At depth d , SELF-QUESTIONING algorithm takes in a visual question Q^d which can be the original visual question ($d = 0$) or a sub-question generated by VQG, a question-related caption C , and hints H^d (if it is available), and try to generate an answer A^d via VQA. If the confidence level of A^d is not high, the SELF-QUESTIONING algorithm starts to raise sub-questions. First, the FQG module takes in Q^d , context C , and hints H^d as input and raises a set of factual questions Q_f . Each question in Q_f is answered by the FQA module and we denote the answer as A_f . Each Q_f and its answer A_f is merged into a factual statement h_f via the QA2H module

and the statement is appended to hints H^d to form H^{d+1} . Second, the VQG module takes in Q^d , context C , and hints H^{d+1} and raises a set of visual questions Q^{d+1} .

B Visualization of Recursive Thinking Process

Figure 10 shows a complete recursive thinking process of our SOCRATIC QUESTIONING method. It involves 4 additional questions to acquire additional information to answer the target question. From this example, we see that LLMs, such as GPT-3 or ChatGPT, have strong capabilities not only in reasoning but also self-questioning. Given the target question to be answered, “*Why are the children wearing hats?*”, LLMs are able to proactively acquire additional commonsense knowledge through factual questions, e.g., “*What are the common reasons why children wear hats?*”, and fine-grained visual information from the input image, e.g., “*What’s the position of the sun in the sky at the time the children are shown wearing hats?*”, “*Are the weather conditions in the image cold or hot?*”. By combining the additional knowledge, e.g., “*cold weather makes people wear hats*” and visual information, e.g., “*it is cold*”, acquired from the recursive Self-Questioning process, the model finally achieves the answer “*warmth*”. This analysis demonstrates that the recursive thinking process of our approach is highly transparent and interpretable.

C Implementation Details

C.1 Language-only Tasks

Implementation Details We leverage ChatGPT (OpenAI, 2022) as the LLM for QA, QG, and QA2H modules. We provide detailed prompts for each module in Appendix K.

Baselines Standard Prompting (SP) prompts ChatGPT to directly answers a question with a few in-context examples. **Chain-of-Thought (CoT)** (Wei et al., 2022) prompts ChatGPT to first generate the thinking process and then generate the answer. We also add the thinking process into the in-context examples. **Self-Consistency Chain-of-Thought (SC-CoT)** (Wang et al., 2023) proposes to run chain-of-thought multiple times on ChatGPT and marginalize the thinking process by taking the most consistent answer. **Tree-of-Thought (ToT)** (Yao et al., 2023) is a recently proposed

framework for improving the reasoning capability of language models. We follow their implementation⁴ which leverages tree-search algorithms to explore the thinking space and select the best thinking path.⁵

Evaluation Metrics For a fair comparison, we use exact match and measure the accuracy for all language-only tasks following previous works (Chowdhery et al., 2023; Hoffmann et al., 2022).

All questions in MMLU Physics, MMLU Chemistry, and LogiQA are multiple-choice questions and the answer is always a single letter like “A”, “B” or “C”. To easily parse the model’s final output, we use “Thus, the final answer is:” as the prefix for the final answers (A or B or C or D, ect.) in the in-context examples for all methods. When we parse the output, we first run a template-based method to extract the answers after “Thus, the final answer is:”. For a few instances (12.52% in CoT, 16.4% in ToT and 11.64% in Socratic Questioning on average) that do not match the template as shown in Figure 4 ToT, the authors manually compare the model’s predictions to the ground truth answers. Thus, we assure that the final performance of all methods is not affected by the output formats.

Datasets **Massive Multitask Language Understanding (MMLU)** (Hendrycks et al., 2020) dataset contains 57 diverse tasks and is used to measure the model’s complex reasoning capability. In this work, we use the physics and chemistry tasks which contain conceptual physics and chemistry multiple-choice questions, respectively. **MATH** (Hendrycks et al., 2021) dataset consists of challenging competition-level mathematics problems which require strong mathematical reasoning ability. **LogiQA** (Liu et al., 2020) dataset contains expert-written questions for testing the logical reasoning capability of humans. For each task, we use the validation set to make design decisions and measure the model’s performance on the test set. The detailed statistics of all datasets can be found in Table 5.

⁴<https://github.com/kyegomez/tree-of-thoughts>

⁵By the time we submit the work, we don’t have access to GPT4 so we use ChatGPT for ToT.

	MATH	MMLU (Physics)	MMLU (Chemistry)	LogiQA
Dev	60	22	26	60
Test	300	235	203	300

Table 5: Statistic of datasets for language-only tasks.

C.2 Multimodal Tasks

Implementation Details We use blip2-flan-t5-xl⁶ as our Visual Perception module. We leverage ChatGPT (OpenAI, 2022) for the FQG, VQG, and FQA modules and GPT-3 (GPT-3-davinci-003) for the VQA module. This decision is motivated by the observation that ChatGPT tends to be excessively cautious and neutral, and avoids answering some questions. We provide detailed sample prompts for each module in Appendix K.

Baselines **BLIP-2** (Li et al., 2023) is a pre-trained vision-language model that leverages an efficient and generic pre-training strategy and is able to follow text prompts. We use the released blip2-flan-t5-xl checkpoint. **PICa** (Yang et al., 2022) prompts GPT-3 with generic image captions to solve VQA in an in-context learning manner. In our experiments, we implement PICa by using blip2-flan-t5-xl as the image captioning model and GPT-3-davinci-003 as the LLM.

Evaluation Metrics We employ the conventional VQA accuracy metric (Goyal et al., 2017) to measure the performance. To alleviate stringent penalization for minor discrepancies between predicted answers and ground truth, we normalize the answers by converting plural forms to singular forms and changing the tense of verbs to present tense. In addition, to address the limitation due to synonyms and expression differences, we employ ChatGPT to evaluate the correctness of the predicted answers (Fu et al., 2023; Liu et al., 2023b). We provide ChatGPT with the visual question, the predicted answer and the ground-truth answer, and ask if the ground-truth answer and the predicted answer can support each other. If the answer is "Yes", we treat the predicted answer as correct. We show the exact prompts used for ChatGPT in Appendix K.8.

Datasets **VQA-V2** (Goyal et al., 2017) is a dataset containing open-ended questions about images. **OK-VQA** (Marino et al., 2019) requires

	VQA-v2	OK-VQA	AOK-VQA
Dev	100	100	100
Test	500	462	444

Table 6: Statistic of datasets for multi-modalities tasks.

model to leverage external knowledge to answer visual questions. **AOK-VQA** (Schwenk et al., 2022) is an augmented successor of OK-VQA, which require commonsense knowledge and strong reasoning capabilities to answer its questions. For each task, we use the validation set to make design decisions and measure the model’s performance on the test set. The detailed statistics of all datasets can be found in Table 6 and Appendix E.

D SELF-QUESTIONING in the Multimodal Setting

See Figure 9.

E Data Leakage in BLIP-2 and GPT-3

Model	VQA-V2	OK-VQA	AOK-VQA
BLIP-2	1.46	2.93	28.08
GPT-3	35.88	23.95	20.4

Table 7: Traditional VQA Accuracy (%) under the setting where no image is provided in the input.

In our preliminary experiments, we discovered an issue that pre-trained models could be subject to data leakage during their pre-training stage. We observed that the baseline models (i.e., BLIP-2 and GPT-3) achieved unjustifiably high performance across all three VQA datasets even without taking images as inputs (see Table 7). To address this issue, we applied a filtering process to remove such contaminated instances. We first test the BLIP-2 and GPT-3 on zero-shot VQA tasks while replacing the original input image with an image composed entirely of black pixels of the same size. Then, we only retain the samples where the models failed to yield a correct answer when the original image is not given. After the filtering, we adopt the 500, 462, and 444 test samples for VQA-V2, OK-VQA, and AOK-VQA, respectively. We use these clean examples for the evaluation throughout the rest of our experiments.

⁶<https://huggingface.co/Salesforce/blip2-flan-t5-xl>




Figure 9: The overview of the SELF-QUESTIONING Algorithm.

F Visualization of Complete SOCRATIC QUESTIONING

See Figure 10.

G Concrete Example

G.1 Large Maximum Number of Turn

Due to the calibration error in LLMs (Jiang et al., 2021), sometimes the pre-trained model’s confidence is not aligned with the answer’s correctness. Thus, in such cases, the model predicts “low” or “medium” confidence in correct answers in the early turns and hence misses the correct answers. If we use fewer turns, we can keep the answer in the early turn regardless of the confidence and hence alleviate the calibration error. Below we show a concrete example in which the model predicts the correct answer in 2 turns and predicts the incorrect answer in 3 turns. When we increase the number of turns, Socratic Questioning may raise some less relevant sub-questions and hence introduce noisy information in the reasoning process. This noisy information can confuse the model, leading to incorrect responses to the original question. For example, consider a simple physics question:

The speed of sound is slightly greater on a ["A. cold day", "B. hot day", "C. day with steady temperature", "D. None of these"]?

In a 2-turn setting, our approach obtains hints: (1) “The speed of sound increases with increasing temperature.”, and (2) “Humidity is a factor in the speed of sound.” According to the hints, it is obvious that the correct answer is B, which is chosen by

our approach in the second turn with the “middle” confidence. In a 3-turn setting, since the LLM does not assign “high” confidence to the answer in the 2 turn, our approach goes deeper in the third turn and gets more information (e.g., (3) “The speed of sound can be affected by several factors, including temperature, humidity and density of the medium.”, (4) “The speed of sound depends on the density and elasticity of the medium it is traveling through, in terms of physical properties.”, (5) “The speed of sound increases with humidity as a result of increased air density.”) As a result, by considering more hints, we potentially introduce less relevant information to the LLM and the noisy information causes the LLM to change its answer to D.

G.2 Large Maximum Number of Depth

We observe that as the depth increases, the context information in the original questions start to vanish and the answers to the sub-questions may be inaccurate in the context of the original question. Thus, by adding the answers to sub-question in larger depth as hints, we can introduce noises to the reasoning process of the LLM which results in wrong answers. Consider a physics question example:

When a spinning system contracts in the absence of an external torque, its rotational speed increases, and its angular momentum [A. decreases, B. increases, C. remains unchanged, D. may increase or decrease]?

Socratic Questioning raises a sub-question: “What affects the rotational speed of a spinning system?” The initial answer to this sub-question

is “*Conservation of angular momentum*”, which provides enough information to answer the original question. In a larger depth setting, the Socratic Questioning raises a deeper sub-question: “*What is the relationship between rotational speed and angular momentum in a spinning system?*” The answer to this question is: “*The angular momentum is directly proportional to the rotational speed*”. Incorporate this hint, the Socratic Questioning changes the answer of the first sub-question to: “*The angular momentum is directly proportional to the rotational speed.*”, which results in an incorrect final answer B.

H Evaluation of Computational Cost

In Table 8, we provide the theoretical number of calls in CoT, SC-CoT, ToT and Socratic Questioning in 2 and 3 turns settings. We also provide the empirical results of the average number of calls per instance and average running time per instance in seconds for all methods. For SC-CoT, we fix the number of calls to 20 times on all the datasets based on the performance curve in (Wang et al., 2023). In ToT, k represents the number of thoughts allowed to be generated per step, T represents the maximum number of steps and b represents the maximum number of states to keep at each step in BFS. Following (Yao et al., 2023), we set k=5, T=3, and b=4. In Socratic Questioning, q represents the maximum number of raised sub-questions for a parent node.

As one can observe, Socratic Questioning with 2 turns and 3 turns achieves better efficiency compared to SC-CoT and ToT. The main reason is that, in the experimental datasets, most questions do not require a large amount of thinking steps to reach the correct answers. Socratic Questioning, adaptively raises sub-questions based on the complexity of the original question and arrives at the correct answer without reaching the theoretical maximum number of turns or depth. In contrast, both SC-COT and ToT employ fixed settings for the number of thoughts generated per step. For relatively straightforward questions, these fixed settings introduce high computational overhead, making the algorithms less efficient in these questions.

I Experimental Results on Other QA and Math Datasets

Table 9 provides the performance of our method and two strong baselines on GSM8K and Strate-

gyQA datasets. As one can observe, our method has significant performance improvement compared to baselines. We use ChatGPT with temperature 0.7 for all methods. For SC-CoT, we sample 20 reasoning paths.

We tried our best to reproduce the results of CoT and SC-CoT reported in (Wang et al., 2023) on StrategyQA. Following (Wang et al., 2022), we use the question-only set from BIG-bench collaboration (2021) and use the exact same prompt template and in-context examples in SC-CoT. However, we cannot reproduce the results on StrategyQA in (Geva et al., 2021) since Code-davinci-002 and Code-davinci-001 are no longer publicly available. In addition, our results of ChatGPT on StrategyQA also agree with more recent studies in (Qin et al., 2023).

J Experiment Results based on GPT-4

To showcase the generalizability of our approach, we have run CoT and Socratic Questioning on MMLU Chemistry and LogiQA based on GPT-4. The experimental results show that our Socratic Questioning approach still significantly outperforms CoT.

K Prmopt Templates

To make our method generalize to other reasoning domains, we carefully design in-context demonstrations to guide the LLM to generate more basic sub-questions in an efficient manner. More concretely, to create high-quality sub-questions in the in-context examples, we take the human reasoning process and domain knowledge into account and carefully annotate the sub-questions by ensuring that they are more basic questions compared to the original question and their solutions can contribute to the reasoning process of the original questions. For examples of sub-questions, please refer to Figure 12. Based on our experiments in math, physics, chemistry and VQA domains, we argue that with a few examples (5 in all our experiments) Socratic-Questioning can generalize to a new domain.

Following (Kadavath et al., 2022), we ask the LLM itself to output a confidence level, “high”, “middle”, or “low”, towards its answer. In the in-context demonstrations, we label the correct answers with supportive hints in the context as “high” confidence, label the correct answers without supportive hints as “middle” confidence, and label incorrect answers as “low” confidence. In this way,

	Standard-Prompting	CoT	SC-CoT	ToT	Socratic Questioning (2 turns)	Socratic Questioning (3 turns)
Theoretical Number of Calls	1	1	20	$k + b*k*(T-1)$	$3 \times \sum_{i=1}^{d-1} [q \times (t-1)]^i$	$3 \times \sum_{i=1}^{d-1} [q \times (t-1)]^i$
Avg. Calls per Instance	1	1	20	31.1	9.22	18.7
Avg. Running Time per Instance (second)	0.33	3.35	67.09	77.99	34.15	53.65

Table 8: Evaluation of computational cost of different methods.

	GSM8K	StrategyQA
CoT	79.0	59.7
SC-CoT	86.0	63.0
Socratic-Questioning	89.33	65.33

Table 9: Accuracy (%) on GSM8K and StrategyQA using Exact Match. The best performance is highlighted in **bold**.

	MMLU Chemistry	LogiQA
CoT	80.2	70.3
Socratic-Questioning	85.73	75.3

Table 10: Accuracy (%) of GPT-4 based approaches using Exact Match. The best performance is highlighted in **bold**.

we can guide the model to align its confidence to the correctness of the predicted answers. Our algorithm will continue raising sub-questions if the estimated confidence is not “high”. Please refer to Figure 11 for more examples.

K.1 Prompts template of QA module

See Figure 11.

K.2 Prompts template of QG module

See Figure 12.

K.3 Prompts template of FQG

See Figure 13.

K.4 Prompts template of FQA

See Figure 14.

K.5 Prompts template of VQG

See Figure 15.

K.6 Prompts template of VQA

See Figure 16 and 17.

K.7 Prompts template of QA-to-Hint

See Figure 18.

K.8 Prompt for chatGPT for semantic-based accuracy evaluation

See Figure 19.



Figure 10: Visualization of a complete recursive thinking process of SOCRATIC QUESTIONING (2-Depth 2-Turn). The process is initialized on the left and is continued on the right.

<p>// System define</p> <p>Imagine you are a thoughtful and logical student. You are given a question. Please use your best judgment to answer the question step by step, and give the confidence. If there are some hints, consider the hints. Note that, the final answer has to be a single letter, which is the ID of the correct option. If there are hints, present which hints you use.</p> <p>// Demonstration</p> <p>Question: The plates of a capacitor are charged to a potential difference of 5 V. If the capacitance is 2 mF, what is the charge on the positive plate?</p> <p>Option: [A. 0.005 C, B. 0.01 C, C. 0.02 C, D. 0.5 C] Note: Form of output is: explanation + [Answer: "option letter"; Used hints: "hints ID or None"; Confidence: "low, middle, or high"].</p> <p>Answer: To find the charge on the positive plate of the capacitor, we can use the formula $Q = CV$, where Q represents the charge, C represents the capacitance, and V represents the potential difference. Given that the potential difference V is 5 V and the capacitance C is 2 mF (millifarads), we can substitute these values into the formula to find the charge Q. $Q = (2 \text{ mF}) * (5 \text{ V})$ To simplify the calculation, we need to convert millifarads (mF) to farads (F). Since 1 mF is equal to 0.001 F, we can rewrite the capacitance as $2 * 0.001 \text{ F}$. $Q = (2 * 0.001 \text{ F}) * (5 \text{ V}) = 0.01 \text{ C}$ Thus, the charge on the positive plate of the capacitor is 0.01 C. Thus, the final answer is: [Answer: B; Used hints: None; Confidence: low]</p> <p>Question: How much work is required to charge a 10 μF capacitor to a potential difference of 100 V?</p> <p>Option: [A. 0.005 J, B. 0.01 J, C. 0.05 J, D. 0.1 J] Note: Form of output is: explanation + [Answer: "option letter"; Used hints: "hints ID or None"; Confidence: "low, middle, or high"].</p> <p>Answer: To calculate the work required to charge a capacitor, we can use the formula: $Work = 0.5 * C * V^2$, where C represents the capacitance and V represents the potential difference. Given that the capacitance C is 10 μF (microfarads) and the potential difference V is 100 V, we can substitute these values into the formula. $Work = 0.5 * (10 \mu\text{F}) * (100 \text{ V})^2 = 5 * 10^{-5} \text{ F} * 10^4 \text{ V}^2 = 0.5 * 10^{-1} \text{ F} * 1 \text{ V}^2 = 5 * 10^{-2} \text{ J}$. Thus, the final answer is: [Answer: C; Used hints: None; Confidence: low]</p>	<p>Question: Ultraviolet light has a wavelength of about $6 \times 10^{-8} \text{ m}$. What is the frequency of this light?</p> <p>Option: [A. $5 \times 10^{15} \text{ Hz}$, B. 0.5 Hz, C. 2 Hz, D. 20 Hz]</p> <p>Hints: (1) Wavelength is inversely proportional to the frequency.</p> <p>Note: Form of output is: explanation + [Answer: "option letter"; Used hints: "hints ID or None"; Confidence: "low, middle, or high"].</p> <p>Answer: To find the frequency of ultraviolet light with a given wavelength, we can use the relationship between wavelength and frequency, which is hint 1: wavelength is inversely proportional to frequency. Mathematically, this relationship can be expressed as: $\text{frequency} = \text{speed of light} / \text{wavelength}$. The speed of light in a vacuum is approximately $3 \times 10^8 \text{ meters per second (m/s)}$. Given that the wavelength of the ultraviolet light is $6 \times 10^{-8} \text{ m}$, we can substitute these values into the formula: $\text{frequency} = (3 \times 10^8 \text{ m/s}) / (6 \times 10^{-8} \text{ m}) = 0.5 \times 10^{16} \text{ Hz} = 5 \times 10^{15} \text{ Hz}$. Thus, the frequency of ultraviolet light with a wavelength of $6 \times 10^{-8} \text{ m}$ is 0.5 Hz. Thus, the final answer is: [Answer: A; Used hints: 1; Confidence: high]</p> <p>Question: A whiffle ball is tossed straight up, reaches the highest point, and falls back down. Air resistance is not negligible. Which of the following statements are true? I. The ball's speed is zero at the highest point. II. The ball's acceleration is zero at the highest point. III. The ball takes a longer time to travel up to the highest point than to fall back down.</p> <p>Option: [A. I only, B. II only, C. I & II only, D. I & III only] Note: Form of output is: explanation + [Answer: "option letter"; Used hints: "hints ID or None"; Confidence: "low, middle, or high"].</p> <p>Answer: Let's evaluate each statement one by one: I. The ball's speed is zero at the highest point. This statement is true. At the highest point of the ball's trajectory, its vertical velocity becomes zero. Since speed is the magnitude of velocity, which is a vector quantity, the speed of the ball indeed becomes zero at the highest point. This is because the ball momentarily stops moving upward before it starts descending. II. The ball's acceleration is zero at the highest point. This statement is false. The ball's acceleration is not zero at the highest point. Even though the ball momentarily stops changing its direction (from upward to downward motion) at the highest point, it still experiences the force of gravity acting downward. The presence of gravity causes the ball's acceleration to remain constant throughout its motion, regardless of the point in its trajectory. III. The ball takes a longer time to travel up to the highest point than to fall back down. This statement is false.</p>	<p>Due to the presence of air resistance, the ball experiences a drag force opposing its motion. As a result, the ball takes longer to reach the highest point of its trajectory compared to the time it takes to fall back down. Air resistance acts as a damping force, reducing the upward speed of the ball and increasing the time it takes to reach the peak. When falling back down, the ball's downward speed is increased by the force of gravity, making the descent faster than the ascent. Based on the evaluations, we can conclude that only Statement I is true. Thus, the final answer is: [Answer: A; Used hints: None; Confidence: middle]</p> <p>Question: An object of volume $2 \times 10^{-3} \text{ m}^3$ and weight 6 N is placed into a tank of water, where it floats. What percentage of the object's volume is above the surface of the water?</p> <p>Option: [A. 12%, B. 30%, C. 60%, D. 70%]</p> <p>Hints: (1) Density of water is 997 kg/m^3; (2) Object density is 306 kg/m^3.</p> <p>Note: Form of output is: explanation + [Answer: "option letter"; Used hints: "hints ID or None"; Confidence: "low, middle, or high"].</p> <p>Answer: The percentage of the object's volume that is above the surface of the water can be calculated using the formula: Percentage above surface = $[(\text{Density of water} - \text{Density of object}) / \text{Density of water}] \times 100$. In this question, the percentage above surface = $[(997 \text{ kg/m}^3 - 306 \text{ kg/m}^3) / 997 \text{ kg/m}^3] \times 100 = (691 \text{ kg/m}^3 / 997 \text{ kg/m}^3) \times 100 = 0.693 \times 100 \approx 70\%$. Thus, the final answer is: [Answer: D; Used hints: 1, 2; Confidence: high]</p>
--	--	--

Figure 11: Prompt template of QA module.

// System define

Imagine you are a thoughtful and logical question-raiser. You are given a physics question. However, the question is too complex or lack of information to answer. You need to raise some questions to decompose the original question into several simpler sub-questions, or to seek additional information that helps you answer the original question. Important notes: do not use pronouns or indefinite pronoun phrases in your generated questions. The raised question has to be the self-contain question, which means including context if it is needed. Each question can only contain one argument. Do not just ask Yes/No questions.

// Demonstration

Question:

An object of volume $2 \times 10^{-3} \text{ m}^3$ and weight 6 N is placed into a tank of water, where it floats. What percentage of the object's volume is above the surface of the water?

Option:

[A. 12%, B. 30%, C. 60%, D. 70%]

Note:

The raised question has to be a self-contain question. Do not use pronouns or indefinite pronoun phrases in the generated questions. Copy context from the original question if needed.

Deep Questions:

1. When an object floats, what function describes the relationship between the object's volume and weight?
2. What is the density of water?
3. An object of volume $2 \times 10^{-3} \text{ m}^3$ and weight 6 N, what is the object's density?

Question:

Compared with the mass of a uranium atom undergoing fission, the combined masses of the products after fission are

Option:

[A. less, B. more, C. the same, D. zero]

Note:

The raised question has to be a self-contain question. Do not use pronouns or indefinite pronoun phrases in the generated questions. Copy context from the original question if needed.

Deep Questions:

1. What causes the change in mass of a particle before and after fission?

Question:

Things that are equivalent according to the equivalence principle are

Option:

[A. space and time, B. a traveling twin and a stay-at-home twin, C. gravity and acceleration, D. mass and energy]

Note:

The raised question has to be a self-contain question. Do not use pronouns or indefinite pronoun phrases in the generated questions. Copy context from the original question if needed.

Deep Questions:

1. What is the equivalence principle?

Question:

Which of these three elements has the most mass per nucleon?

Option:

[A. Hydrogen, B. Iron, C. Uranium, D. Same in each]

Note:

The raised question has to be a self-contain question. Do not use pronouns or indefinite pronoun phrases in the generated questions. Copy context from the original question if needed.

Deep Questions:

1. What is the nucleon mass of hydrogen?
2. What is the nucleon mass of iron?
3. What is the nucleon mass of uranium?

Question:

A microwave oven is connected to an outlet, 120000 mV, and draws a current of 2 amps. At what rate is energy being used by the microwave oven?

Option:

[A. 10 W, B. 30 W, C. 60 W, D. 240 W]

Note: The raised question has to be a self-contain question. Do not use pronouns or indefinite pronoun phrases in the generated questions. Copy context from the original question if needed.

Deep Questions:

1. Given a given voltage and current, how to calculate the power?
2. How many volts equal 12000 microvolts?

Figure 12: Prompt template of QG module.

// System define

Imagine you are a blind but intelligent asker only given a question and description of an image. You need to ask me (at most 5) questions about facts or commonsense knowledge to help you get more information about the image to help answer the given question. Important notes: do not use pronouns in your generated questions. Each question can only contain one argument. Do not just ask Yes/No questions.

// Demonstration

Image Caption:

- A bowl of oranges in a bowl.

Question:

- What states are these grown in?

Factual Questions:

1. Reason: Each US state has different specialties.
Question: In which state in the USA are oranges grown?
2. Reason: Oranges grow in states where the environment is good for them.
Question: What environments are suitable for orange growth?

// Input

Image Caption:

- A bedroom with a bed and a canopy.

Question:

- Is this a room for a boy or a girl?

Factual Questions:

Figure 13: Prompt template of FQG module.

// System define

Imagine you are a polymath familiar with encyclopedias and all kinds of common-sense knowledge. You need to answer a question about some facts or some common-sense knowledge in short sentence.

// Demonstration

Question:

- What is human life expectancy in the United States?

Answer:

- Human life expectancy in the United States is 78 years.

// Input

Question:

- In which state in the USA are oranges grown?

Answer:

Figure 14: Prompt template of FQA module.

// System define

Imagine you are a blind but intelligent system. You are given the context of an image and a question about the image. However, the current context is insufficient to answer the question. You should ask me at least two short questions about visual information in the image to help you answer the question. Important notes: do not use pronouns in your generated questions. Each question can only contain one argument. Do not just ask Yes/No questions.

// Demonstration

Image Caption:

- Two women walking on a sidewalk with an umbrella.

Question:

- Are the ladies friends?

Visual Prompts:

1. Reason: People with close relationships, such as friends, walk closer.
Prompt: Are the two women walking close to each other?
2. Reason: The body language between friends will be more intimate, such as hugging, holding hands, etc.
Prompt: What's the body language of the two women?

Image Caption:

- A horse pulling a carriage with two people in it.

Hints:

1. People generally use tools like bridles to force horses to work.

Question:

- Does the horse do this because it wants to?

Visual Prompts:

1. Reason: When animals are forced to work, they show facial expressions such as anger and sadness.
Prompt: What is the expression on the horse's face while it's pulling the carriage?
2. Reason: Humans often use tools such as bridles to control animals and force them to work.
Prompt: What type of tools or equipment is being used to control the horse while it pulls the carriage?

// Input

Image Caption:

- A bedroom with a bed and a canopy.

Question:

- Is this a room for a boy or a girl?

Hints:

1. The commonly used colors for boys' interior decoration are blue and gray.
2. The commonly used colors for girls' interior decoration are pink and white.

Visual Prompts:

Figure 15: Prompt template of VQG module.

// System define

Imagine you are a blind but intelligent question answering system. You are asked a visual question about an image. I will provide you the caption of the image and some useful visual hints. Please use your best judgement to answer the visual question.

// Demonstration

Image Caption:

- A man holding a dog on his back.

Hints:

1. Dogs usually use mouth to catch objects
2. The popular game people play with dog is frisbee
3. The man is holding a frisbee

Question:

- Which part of this animal would be in use if it was playing the game that is played with the items the man is holding? (If the information is not enough to answer the question, answer "lack of information")

Answer:

- Hints 1,2,3 are useful. The answer is: mouth

Image Caption:

- A busy city street with many people walking around.

Question:

- Why might someone go to this place? (If the information is not enough to answer the question, answer "lack of information")

Answer:

- Shop

Image Caption:

- A bowl of oranges in a bowl.

Question:

- What states are these grown in? (If the information is not enough to answer the question, answer "lack of information")

Answer:

- Lack of information

// Input

Image Caption:

- A bathroom with a toilet and a sink.

Hint:

1. Toilet could be used by both man and woman
2. There is a razor near the sink

Question:

- Who leaves a toilet like this? (If the information is not enough to answer the question, answer "lack of information")

Answer:

// System define

Imagine you are a blind but intelligent question answering system. You are asked a visual question about an image. I will provide you the caption of the image and some useful visual hints. Please use your best judgement to answer the visual question.

// Demonstration

Image Caption:

- A man holding a dog on his back.

Hints:

1. Dogs usually use mouth to catch objects
2. The popular game people play with dog is frisbee
3. The man is holding a frisbee

Question:

- Which part of this animal would be in use if it was playing the game that is played with the items the man is holding? (Must return an answer. The final answer should be 1 or 2 words (maximum 2 words). If you are not sure, you can guess the most plausible answer)

Answer:

- Hints 1,2,3 are useful. The answer is: mouth

Image Caption:

- A busy city street with many people walking around.

Question:

- Why might someone go to this place? (Must return an answer. The final answer should be 1 or 2 words (maximum 2 words). If you are not sure, you can guess the most plausible answer)

Answer:

- Shop

// Input

Image Caption:

- A bathroom with a toilet and a sink.

Hint:

1. Toilet could be used by both man and woman
2. There is a razor near the sink

Question:

- Who leaves a toilet like this? (Must return an answer. The final answer should be 1 or 2 words (maximum 2 words). If you are not sure, you can guess the most plausible answer)

Answer:

Figure 17: Prompt template of VQA module (force answer).

Figure 16: Prompt template of VQA module.

```
// System define
You are given a question-and-answer pair, can you help
me to merge the question and answer into a statement
sentence. If the question or the answer is ambiguous you
can just output the token "unknown". If the merged
sentence is ambiguous, you can just output the token
"unknown". If you can merge the question-and-answer
pair, just output the sentence.

// Input
Question:
• In which state in the USA are oranges grown?
Answer:
• California
Statement:
```

Figure 18: Prompt template of QA-to-Hint module.

```
// System define
Imagine you are a strict marking teacher or grader. I will
give you a question, a correct answer, and a student
answer. You need to tell me "1" or "0" (where 1 means
correct, 0 means incorrect). "1" does not mean the
student's answer must exactly match the correct answer.
If they have the same meaning for the given question,
then it is also "1". However, an ambiguous answer is "0"
(e.g., correct answer: "1789", student answer: "long long
ago").

// Input
Question:
• In which state in the USA are oranges grown?
Correct Answer:
• California
Student Answer:
• California state
Grade:
```

Figure 19: Prompt template for chatGPT of semantic-based evaluation.